

# ECS 150 - System Calls

## Assignment Project Exam Help

*Joël Porquet*  
<https://powcoder.com>

UC Davis - Spring Quarter 2017  
Add WeChat powcoder

**Readings:** *OSPP Chap 3*

# System calls

## Example

```
#include <unistd.h>
```

```
int main(int argc, char **argv)
```

```
{
```

```
    int fd, nread;
```

```
    char buf[1024];
```

```
    fd = open("/path/to/myfile", O_RDONLY);
```

```
    nread = read(fd, buf, 1024);
```

```
    ...
```

```
    close(fd);
```

```
    return 0;
```

```
}
```

Assignment Project Exam Help

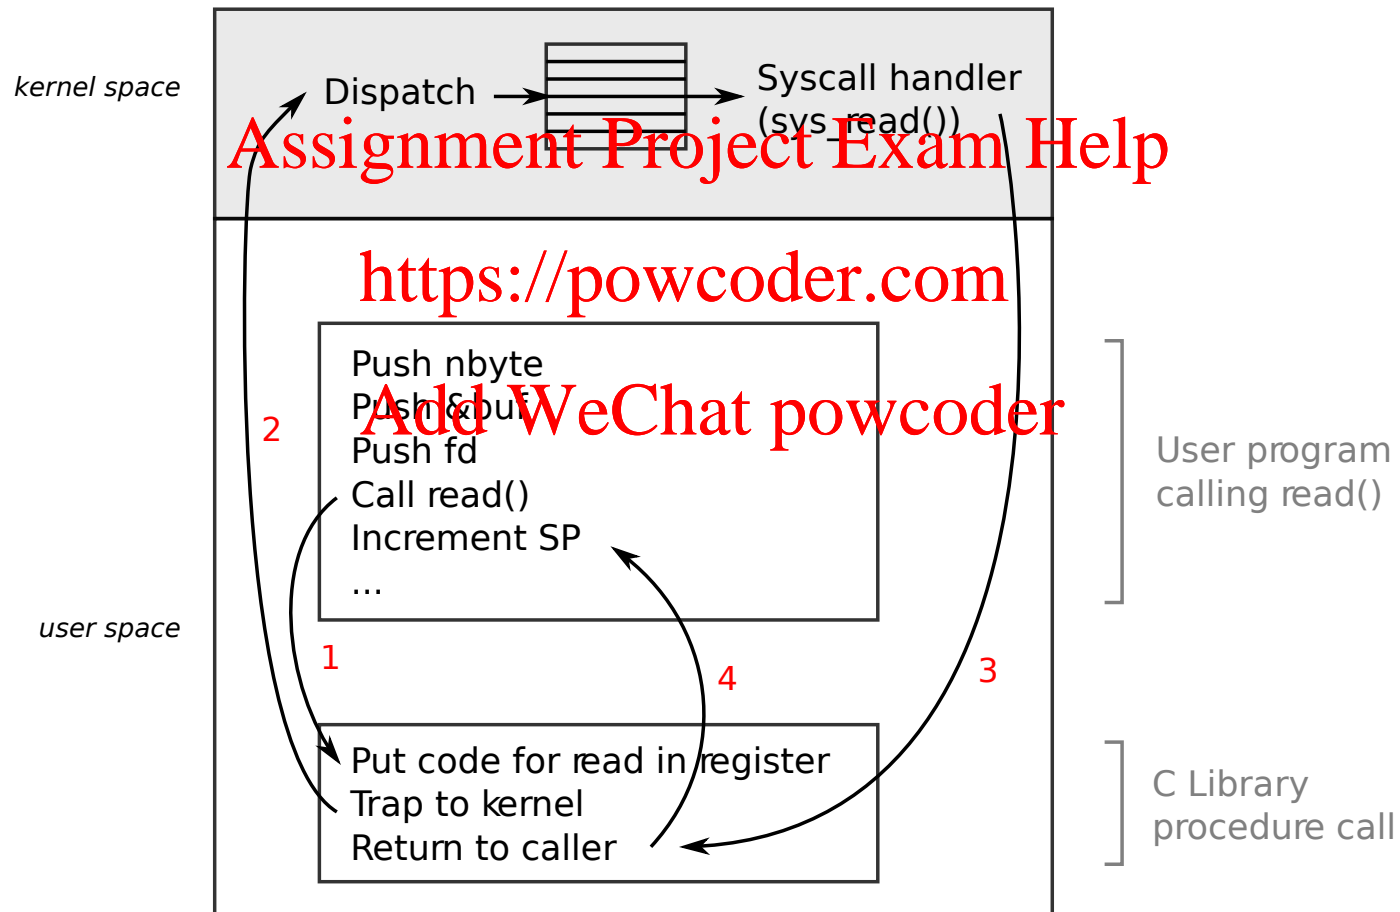
<https://powcoder.com>

Add WeChat powcoder

# System calls

## Overview

- API to the kernel
- Hide kernel and hardware complexity from applications



# System calls

## Inter-communication

How the system calls communicate back to applications?

### Return value

- Usually -1 on error,  $\geq 0$  on success
- C library functions set global variable `errno` to encode the error
  - E2BIG, EACCESS, EAGAIN, EBADF, ENOMEM, ... (man `errno`)
  - Use `perror()` to display a string

```
void *ptr = malloc(1);
if (!ptr) {
    int errno_save = errno;
    perror("malloc");
    fprintf(stderr, "malloc: %s\n", strerror(errno_save));
    exit(EXIT_FAILURE);
}
```

### Buffers

- Part of system calls arguments
- Values need to be copied between user and kernel space

# System calls

Process

Files

Pipe

Signals

Memory

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Process

## Definition

- *A process is a program in execution*
- Each process has its own memory space and process control block
- **Process control block**
  - Kernel structure
  - Stores all information associated with a process
  - Register values for context switching, open files, user ID, group ID, etc.
- Processes are indexed by the process ID (PID)

Assignment Project Exam Help

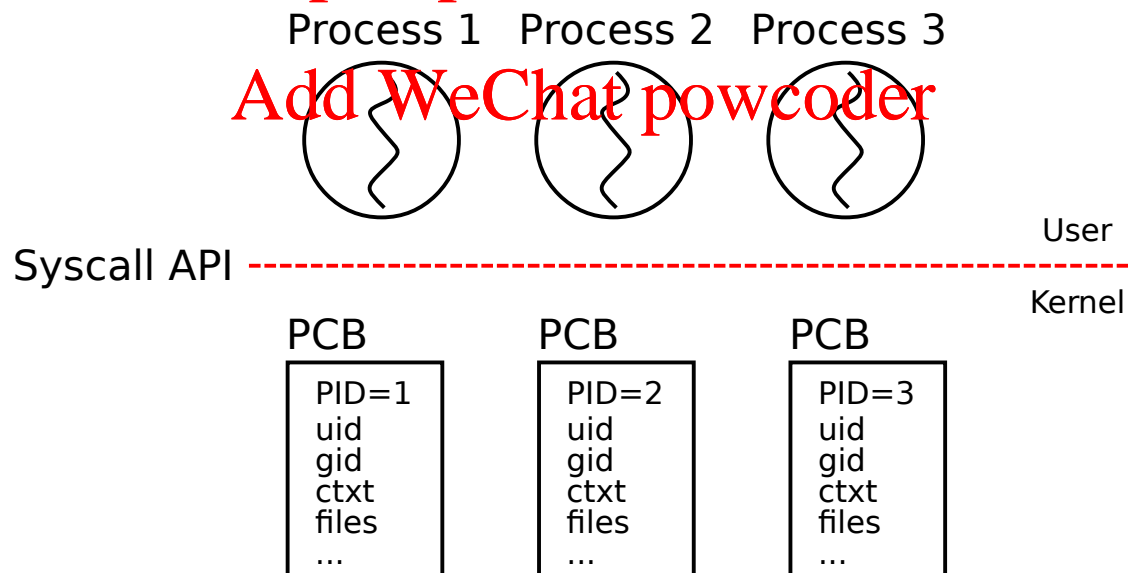
<https://powcoder.com>

Add WeChat powcoder

# Process

## Definition

- *A process is a program in execution*
- Each process has its own memory space and process control block
- **Process control block**
  - Kernel structure
  - Stores all information associated with a process
  - Register values for context switching, open files, user ID, group ID, etc.
- Processes are indexed by the process ID (PID)



# Process

## Related system calls

- `fork()`: Create a new process
- `exec()`: Change executed program in process
- `exit()`: End process
- `wait()/waitpid()`: Wait for a child process and collect exit code
- `getpid()`: Get process PID
- `getpgrp()`: Get process GID

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Process

## fork()

- Syntax: `pid = fork()`
- The child gets almost identical copy of the parent
- File descriptors, arguments, memory, stack, etc. are all copied
- Even the program counter
- Only one thing differs.. Which one?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Process

## Forgetful forking

```
int main(int argc, char **argv)
{
    fork();
    printf("I am the process!\n");
    return 0;
}
```

What gets printed? **Assignment Project Exam Help**

**<https://powcoder.com>**

**Add WeChat powcoder**

# Process

## Forgetful forking

```
int main(int argc, char **argv)
{
    fork();
    printf("I am the process!\n");
    return 0;
}
```

What gets printed? **Assignment Project Exam Help**

```
$ ./simple-fork
I am the process!
I am the process!
```

**<https://powcoder.com>**

**Add WeChat powcoder**

# Process

## Forgetful forking

```
int main(int argc, char **argv)
{
    fork();
    printf("I am the process!\n");
    return 0;
}
```

What gets printed? **Assignment Project Exam Help**

```
$ ./simple-fork
I am the process!
I am the process!
```

<https://powcoder.com>

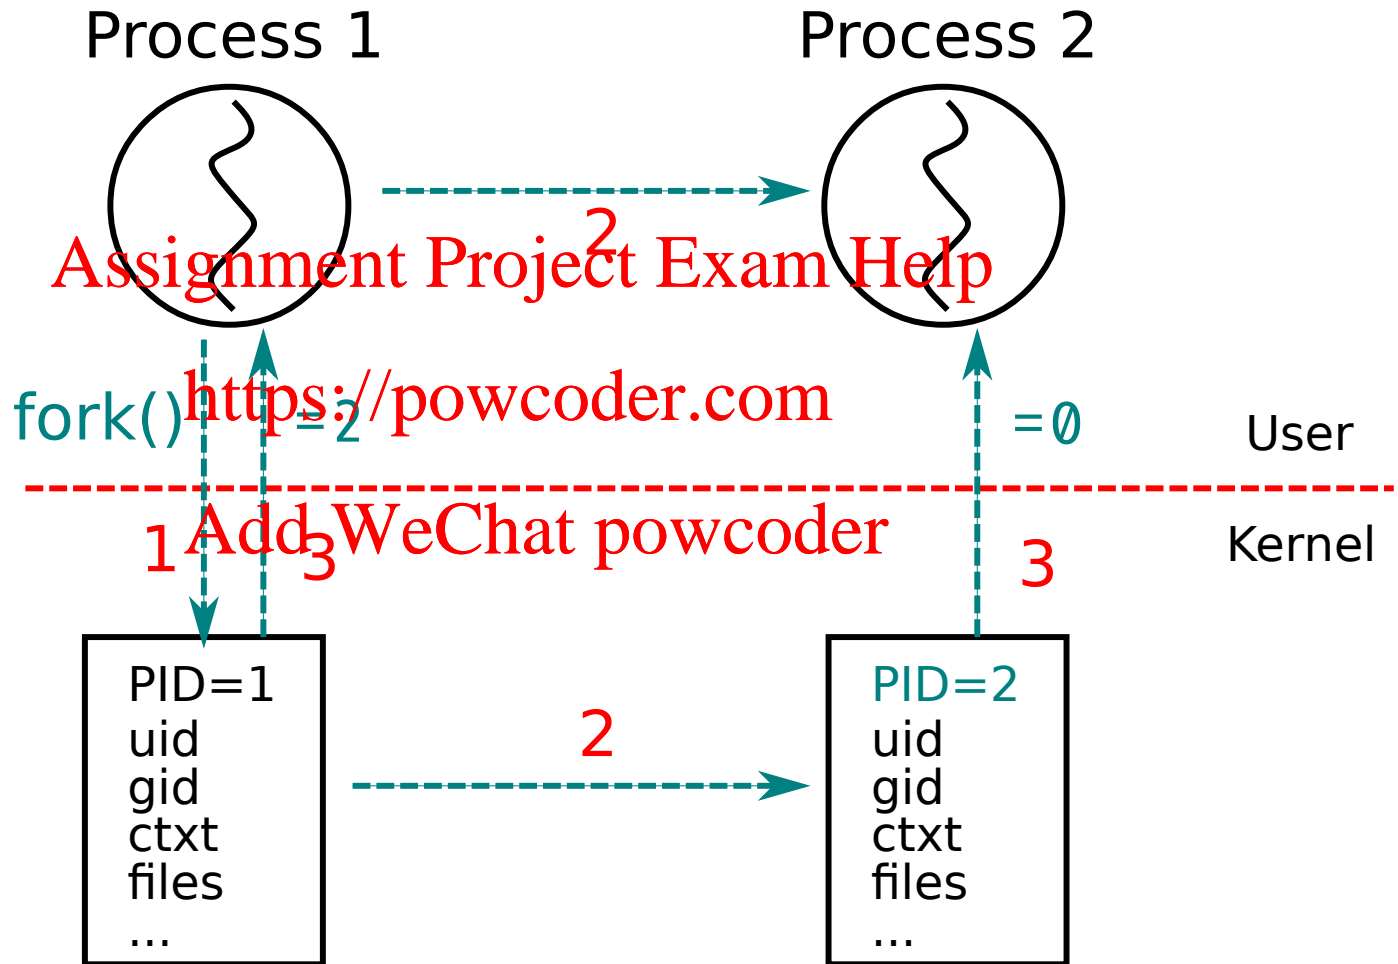
## Using the return value

**Add WeChat powcoder**

fork() returns:

- **zero** for the child
- **PID of the child** for the parent
- **(-1** in case of error)

# Fork illustrated



# Process

## Fork example

```
int main(int argc, char **argv)
{
    pid_t pid;
    pid = fork();
    if (pid > 0)
        printf("I am the parent!\n");
    else if (pid == 0)
        printf("I am the child!\n");
    else
        printf("I am the initial process! But something went wrong...");
    printf("I am here now!");
    return 0;
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

What gets printed? (assuming everything goes fine)

# Process

## Fork example

```
int main(int argc, char **argv)
{
    pid_t pid;
    pid = fork();
    if (pid > 0)
        printf("I am the parent!\n");
    else if (pid == 0)
        printf("I am the child!\n");
    else
        printf("I am the initial process! But something went wrong...");
    printf("I am here now!");
    return 0;
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


What gets printed? (assuming everything goes fine)

```
$ ./complete-fork
I am the parent!
I am here now!
I am the child!
I am here now!
```

# Process

## Process hierarchy

- Notion of a hierarchy (tree) of processes
- Each process has a single parent
- In Unix, all user processes have *init* as their ultimate ancestor



Additional ways to group processes:

- Process groups (signalling)
- Sessions (job control)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Process

## exec()

- Change executed program in current process
- Several different forms with slightly different syntax:
  - `exec[lv]p?e?()` (see man page for details)

```
void exec_ls(void)
{
    char *cmd = "/bin/ls";
    char *args[] = { cmd, ".", NULL };
    char *env[] = { NULL };

    status = execve("/bin/ls", args, env);

    printf("Did everything go well?\n");
    printf("status=%d\n", status);
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Will the `printf()`'s be executed?

# Process

## wait()/waitpid()

- When a process is done, it can either explicitly call `exit(status)` or just return and the exit call will be done implicitly

```
int retval = main(argc, argv);  
exit(retval);
```

- The status is what `$?` reflects in the shell
- A parent can wait for its children (and by default blocks until they are done)

```
int statloc;  
waitpid(pid, &statloc, options);  
wait(&statloc);
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Process

## wait() example

```
pid = fork();  
if (pid == 0) {  
    /* child */  
    printf("I'm the child and I will die soon!\n");  
    exit(42);  
} else {  
    /* parent */  
    int status;  
    wait(&status); /* could also be waitpid(pid, &status, 0) */  
    printf("Child exited with return code %d\n", WEXITSTATUS(status));  
}
```

Assignment Project Exam Help

<https://powcoder.com>

What gets printed?

Add WeChat powcoder

# Process

## wait() example

```
pid = fork();
if (pid == 0) {
    /* child */
    printf("I'm the child and I will die soon!\n");
    exit(42);
} else {
    /* parent */
    int status;
    wait(&status); /* could also be waitpid(pid, &status, 0) */
    printf("Child exited with return code %d\n", WEXITSTATUS(status));
}
```

Assignment Project Exam Help  
<https://powcoder.com>

What gets printed?

```
$ ./fork-wait
I am the child and I will die soon!
Child exited with return code 42
```

Add WeChat powcoder

# Process

## Putting it together: fork()+exec()+wait()

```
int main(int argc, char **argv)
{
    pid_t pid;
    int status;
    char *cmd[3] = { "ls", ".", NULL, };

    pid = fork();
    if (pid == 0) {
        /* Child */
        execvp(cmd[0], cmd);
        perror("execvp");
        exit(1);
    } else if (pid > 0) {
        /* Parent */
        waitpid(-1, &status, 0);
        printf("Child exited with status: %d\n", WEXITSTATUS(status));
    } else {
        perror("fork");
        exit(1);
    }
    return 0;
}
```

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

# Process

## Full example: the shell

- A **Shell** is a programs that typically make heavy use of process system calls
- Basic cycle:
  1. Display prompt
  2. Read line from input
  3. Parse line
  4. Fork
    1. Child executes the command
    2. Parent waits
- Handle background jobs (&)
- Handle redirections (< and >) and pipes (|) for connecting `stdin` or `stdout` of the child to files or other processes

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Process

## Shell skeleton

```
while(1) {                                /* Repeat forever */
    char **command;

    display_prompt();                      /* Display prompt in terminal */
    read_command(&command);               /* Read input from terminal */

    if (fork() != 0) {                    /* fork off child process */
        /* Parent */
        waitpid(-1, &status, 0);          /* wait for child to exit */
    } else {
        execv(command[0], command);       /* execute command */
        perror("execv");                  /* coming back here is an error */
        exit(1);
    }
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder